

CS35L – Winter 2019

Slide set:	6.2
Slide topics:	Multithreaded Programming
Assignment:	6

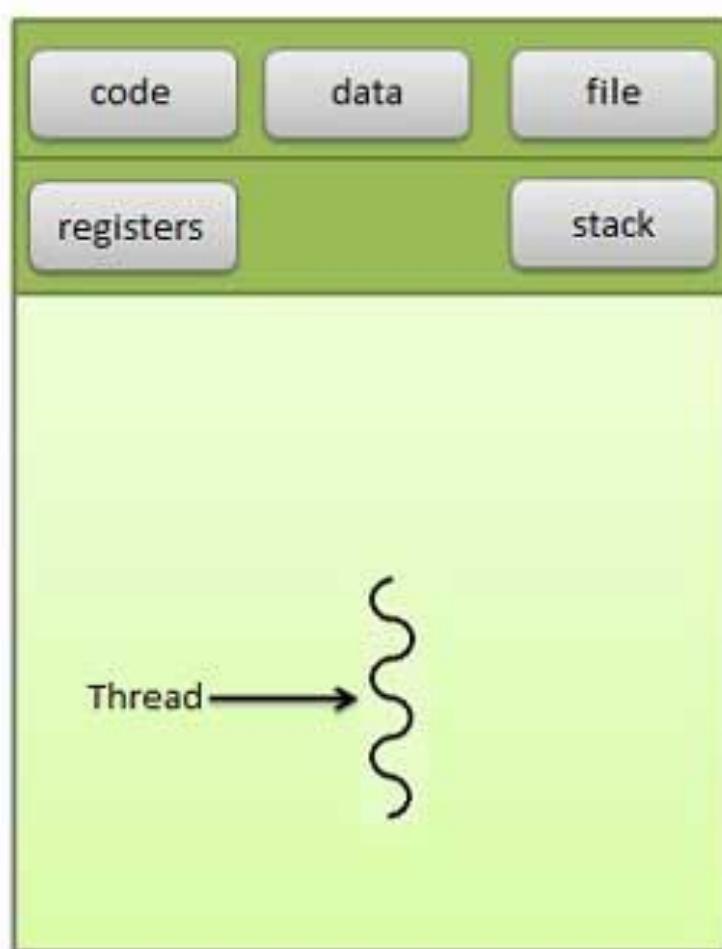
Multitasking

- \$ tr 'abc' 'xyz' | sort -u | comm -23 file1 -
 - Process 1 (tr)
 - Process 2 (sort)
 - Process 3 (comm)
- Each process has its own address space
- How do these processes communicate?
 - Pipes/System Calls

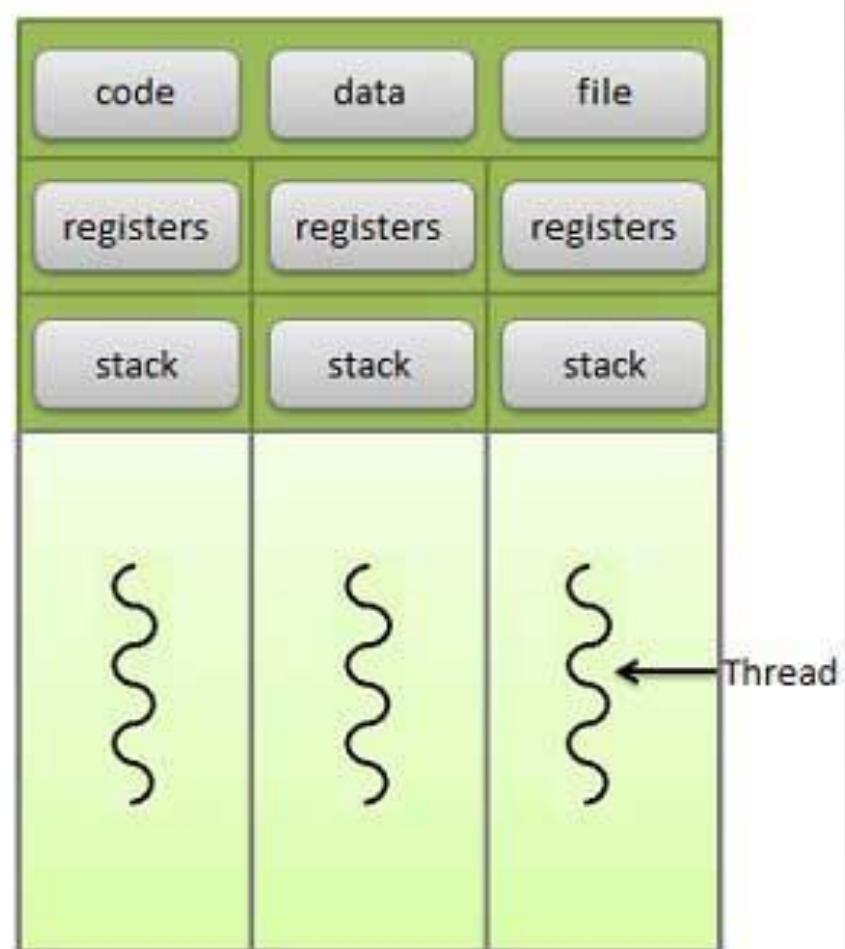
Multithreading

- Threads share all of the process's memory except for their stacks
- => Data sharing requires no extra work (no system calls, pipes, etc.)

Multithreading Memory Layout



Single threaded Process



Multi-threaded Process

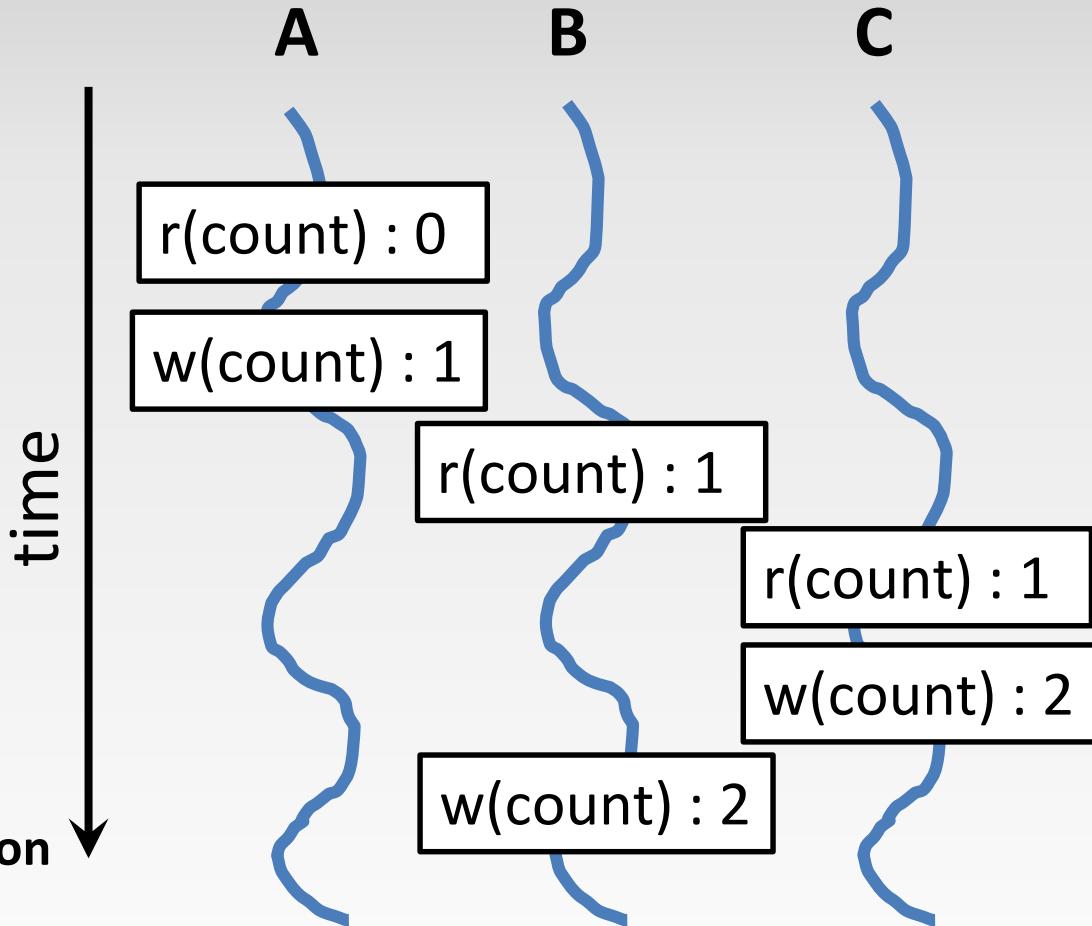
Shared Memory

- Makes multithreaded programming
 - **Powerful**
 - can easily access data and share it among threads
 - **More efficient**
 - No need for system calls when sharing data
 - Thread creation and destruction less expensive than process creation and destruction
 - **Non-trivial**
 - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```

Result depends on order of execution
=> Synchronization needed



Multithreading & Multitasking: Comparison

- **Multithreading**
 - Threads share the same address space
 - Light-weight creation/destruction
 - Easy inter-thread communication
 - An error in one thread can bring down all threads in process
- **Multitasking**
 - Processes are insulated from each other
 - Expensive creation/destruction
 - Expensive IPC (interprocess communication)
 - An error in one process cannot bring down another process

Lab 6

- Evaluate the performance of multithreaded sort
- Add /usr/local/cs/bin to PATH
 - \$ export PATH=/usr/local/cs/bin:\$PATH
- Generate a file containing 10M random single-precision floating point numbers, one per line with no white space
 - /dev/urandom: pseudo-random number generatorDisk quota exceeded
 - <http://www.seasnet.ucla.edu/seasnet-account-quotas/>

Lab 6

- od
 - write the contents of its input files to standard output in a user-specified format
 - Options
 - -t : select output format
 - -N <count>: Format no more than *count* bytes of input
- sed, tr
 - Remove address, delete spaces, add newlines between each float

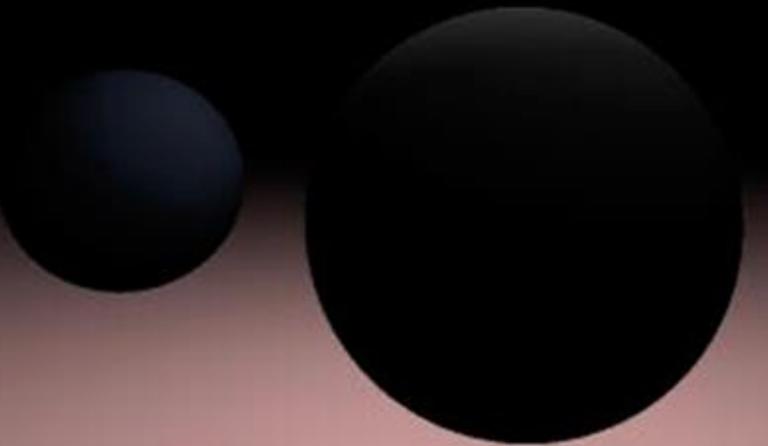
Lab 6

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the `-g` option: compare by general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g file_name > /dev/null (1 thread)`
 - `$ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null`
 - Record the times and steps in `log.txt`

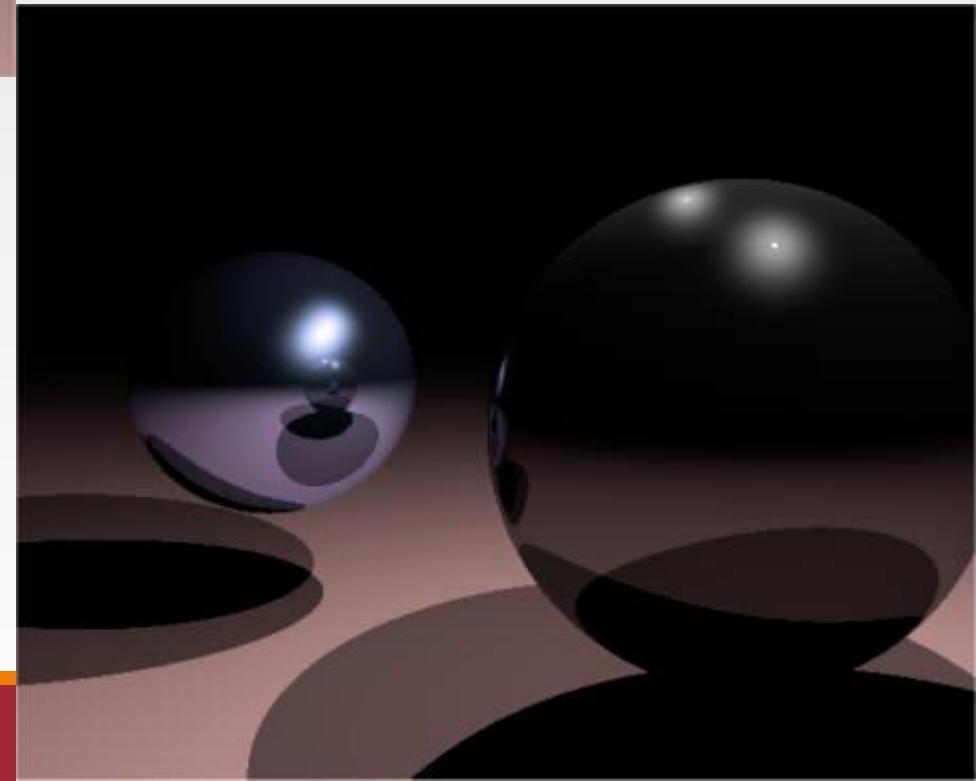
Ray Tracing

- An advanced computer graphics technique for rendering 3D images
- Mimics the propagation of light through objects
- Simulates the effects of a single light ray as it's reflected or absorbed by objects in the images

With ray tracing



Without ray tracing



Computational Resources

- Ray Tracing produces a very high degree of visual realism at a high cost
 - The algorithm is *computationally intensive*
- => Good candidate for multithreading
(embarrassingly parallel)

Homework 6

- Download the single-threaded ray tracer implementation
- Run it to get output image
- Multithread ray tracing
- Run the multithreaded version and compare resulting image with single-threaded one

Basic pthread Functions

include `<pthread.h>` and link with the `-lpthread` library

There are 5 basic pthread functions:

- 1. `pthread_create`:** creates a new thread within a process
- 2. `pthread_join`:** waits for another thread to terminate
- 3. `pthread_equal`:** compares thread ids to see if they refer to the same thread
- 4. `pthread_self`:** returns the id of the calling thread
- 5. `pthread_exit`:** terminates the currently running thread

`pthread_create`

- **Function:** creates a new thread and makes it executable
- Can be called any number of times from anywhere within code
- Return value:
 - Success: zero
 - Failure: error number

Parameters

```
int pthread_create( pthread_t *tid, const pthread_attr_t *attr,  
                    void *(my_function)(void *), void *arg );
```

- **tid**: unique identifier for newly created thread
- **attr**: object that holds thread attributes (priority, stack size, etc.)
 - Pass in NULL for default attributes
- **my_function**: function that thread will execute once it is created
- **arg**: a *single* argument that may be passed to **my_function**
 - Pass in NULL if no arguments

pthread_create Example

```
#include <pthread.h> ...
void *printMsg(void *thread_num) {
    int t_num = (int) thread_num;
    printf("It's me, thread #%d!\n", t_num); }

int main() {
    pthread_t tids[3];
    int t;
    for(t = 0; t < 3; t++) {
        ret = pthread_create(&tids[t], NULL, printMsg, (void *) t);
        if(ret) {
            printf("Error creating thread. Error code is %d\n", ret");
            exit(-1); }
    }
}
```

Possible problem with this code?

If main thread finishes before all threads finish their job -> incorrect results

`pthread_join`

- **Function:** makes originating thread wait for the completion of all its spawned threads' tasks
- Without join, the originating thread would exit as soon as it completed its job
 - ⇒ A spawned thread can get aborted even if it is in the middle of its chore
- Return value:
 - ⇒ Success: zero
 - ⇒ Failure: error number

Arguments

```
int pthread_join(pthread_t tid, void **status);
```

- **tid:** thread ID of thread to wait on
- **status:** the exit status of the target thread
is stored in the location pointed to by
`*status`
 - Pass in `NULL` if no status is needed

pthread_join Example

```
#include <pthread.h> ...
#define NUM_THREADS 5

void *PrintHello(void *thread_num) {
    printf("\n%d: Hello World!\n", (int) thread_num); }

int main() {
    pthread_t threads[NUM_THREADS];
    int ret, t;
    for(t = 0; t < NUM_THREADS; t++) {
        printf("Creating thread %d\n", t);
        ret = pthread_create(&threads[t], NULL, PrintHello, (void *) t);
        // check return value for errors }

    for(t = 0; t < NUM_THREADS; t++) {
        ret = pthread_join(threads[t], NULL);
        // check return value for errors }
}
```

Homework 6

- Build a multi-threaded version of Ray tracer
- Modify “main.c” & “Makefile”
 - Include <pthread.h> in “main.c”
 - Use “pthread_create” & “pthread_join” in “main.c”
 - Link with -lpthread flag (LDLIBS target in Makefile)
- make clean check
 - Outputs “1-test.ppm”
 - Can see “1-test.ppm” in GIMP/Image viewer

baseline.ppm & 1-test.ppm

